

AI

Artificial Intelligence

8.1.3

Agent-based view of AI (Ch. 2)

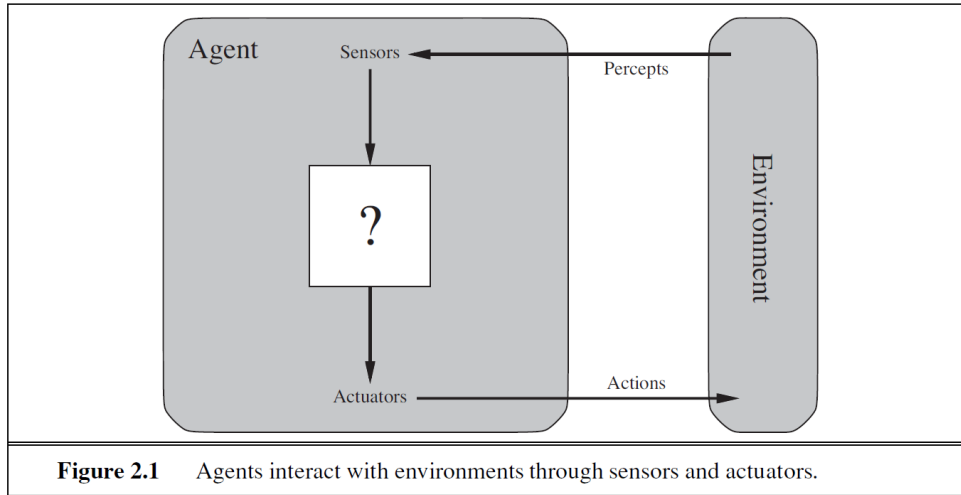
Outline

- Agents and environments
- Rationality
- PEAS (Performance measure, Environment, Actuators, Sensors)
- Environment types
- Agent types

Agents


- An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**
- Human agent: eyes, ears, and other organs for sensors; hands, legs, mouth, and other body parts for actuators
- Robotic agent: cameras and infrared range finders for sensors; various motors for actuators

Agents and environments



- The **agent function** maps from percept histories to actions:
 $[f: \mathcal{P}^* \rightarrow \mathcal{A}]$
- The **agent program** runs on the physical **architecture** to produce f
- agent = architecture + program

Vacuum-cleaner world



- Percepts: location and contents, e.g., [A,Dirty]
- Actions: *Left*, *Right*, *Suck*, *NoOp*

A vacuum-cleaner agent

Percept sequence	Action
$[A, \textit{Clean}]$	\textit{Right}
$[A, \textit{Dirty}]$	\textit{Suck}
$[B, \textit{Clean}]$	\textit{Left}
$[B, \textit{Dirty}]$	\textit{Suck}
$[A, \textit{Clean}], [A, \textit{Clean}]$	\textit{Right}
$[A, \textit{Clean}], [A, \textit{Dirty}]$	\textit{Suck}
\vdots	\vdots

function REFLEX-VACUUM-AGENT($[location, status]$) **returns** an action

if $status = \textit{Dirty}$ **then return** \textit{Suck}

else if $location = A$ **then return** \textit{Right}

else if $location = B$ **then return** \textit{Left}

What is the **right** function?

Can it be implemented in a small agent program?

Rational agents

- An agent should strive to "do the right thing", based on what it can perceive and the actions it can perform. The right action is the one that will cause the agent to be most successful
- Performance measure: An objective criterion for success of an agent's behavior
- E.g., performance measure of a vacuum-cleaner agent could be amount of dirt cleaned up, amount of time taken, amount of electricity consumed, amount of noise generated, etc.

Rational agents

- **Rational Agent:** For each possible percept sequence, a rational agent should select an action that is expected to maximize its performance measure, given the evidence provided by the percept sequence and whatever built-in knowledge the agent has.

Rational agents

- Rationality is distinct from omniscience (all-knowing with infinite knowledge)
- Agents can perform actions in order to modify future percepts so as to obtain useful information (information gathering, exploration)
- An agent is **autonomous** if its behavior is determined by its own experience (with ability to learn and adapt)

PEAS

- PEAS
 - Performance measure, Environment, Actuators, Sensors
- Must first specify the setting for intelligent agent design
- Consider, e.g., the task of designing an automated taxi driver:
 - Performance measure:
 - Safe, fast, legal, comfortable trip, maximize profits
 - Environment:
 - Roads, other traffic, pedestrians, customers
 - Actuators:
 - Steering wheel, accelerator, brake, signal, horn
 - Sensors:
 - Cameras, sonar, speedometer, GPS, odometer, engine sensors, keyboards

PEAS

- Agent: Medical diagnosis system
 - Performance measure:
 - Healthy patient, minimize costs, lawsuits
 - Environment:
 - Patient, hospital, staff
 - Actuators:
 - Screen display (questions, tests, diagnoses, treatments, referrals)
 - Sensors:
 - Keyboard (entry of symptoms, findings, patient's answers)

PEAS

- Agent: Part-picking robot
 - Performance measure:
 - Percentage of parts in correct bins
 - Environment:
 - Conveyor belt with parts, bins
 - Actuators:
 - Jointed arm and hand
 - Sensors:
 - Camera, joint angle sensors

PEAS

- Agent: Interactive English tutor
 - Performance measure:
 - Maximize student's score on test
 - Environment:
 - Set of students
 - Actuators:
 - Screen display (exercises, suggestions, corrections)
 - Sensors:
 - Keyboard

PEAS

- Agent: [let's think of another one]
 - Performance measure:
 - Environment:
 - Actuators:
 - Sensors:

Internet shopping agent

Performance measure?? price, quality, appropriateness, efficiency

Environment?? current and future WWW sites, vendors, shippers

Actuators?? display to user, follow URL, fill in form

Sensors?? HTML pages (text, graphics, scripts)

Environment types

- **Fully observable** (vs. partially observable): An agent's sensors give it access to the complete state of the environment at each point in time.
- **Deterministic** (vs. stochastic): The next state of the environment is completely determined by the current state and the action executed by the agent. (If the environment is deterministic except for the actions of other agents, then the environment is **strategic**)
- **Episodic** (vs. sequential): The agent's experience is divided into atomic "episodes" (each episode consists of the agent perceiving and then performing a single action), and the choice of action in each episode depends only on the episode itself.

Environment types

- **Static** (vs. dynamic): The environment is unchanged while an agent is deliberating. (The environment is **semidynamic** if the environment itself does not change with the passage of time but the agent's performance score does)
- **Discrete** (vs. continuous): A limited number of distinct, clearly defined percepts and actions.
- **Single agent** (vs. multiagent): An agent operating by itself in an environment.

Environment types

Chess with a clock	Chess without a clock	Taxi driving
-----------------------	--------------------------	--------------

Fully observable

Deterministic

Episodic

Static

Discrete

Single agent

Environment types

	Chess with a clock	Chess without a clock	Taxi driving
Fully observable	Yes	Yes	No
Deterministic	Strategic	Strategic	No
Episodic	No	No	No
Static	Semidynamic	Yes	No
Discrete	Yes	Yes	No
Single agent	No	No	No

- The environment type largely determines the agent design
- The real world is (of course) partially observable, stochastic, sequential, dynamic, continuous, multi-agent

Agent functions and programs

- An agent is completely specified by the agent function mapping percept sequences to actions
- One agent function (or a small equivalence class) is rational
- Aim: find a way to implement the rational agent function concisely

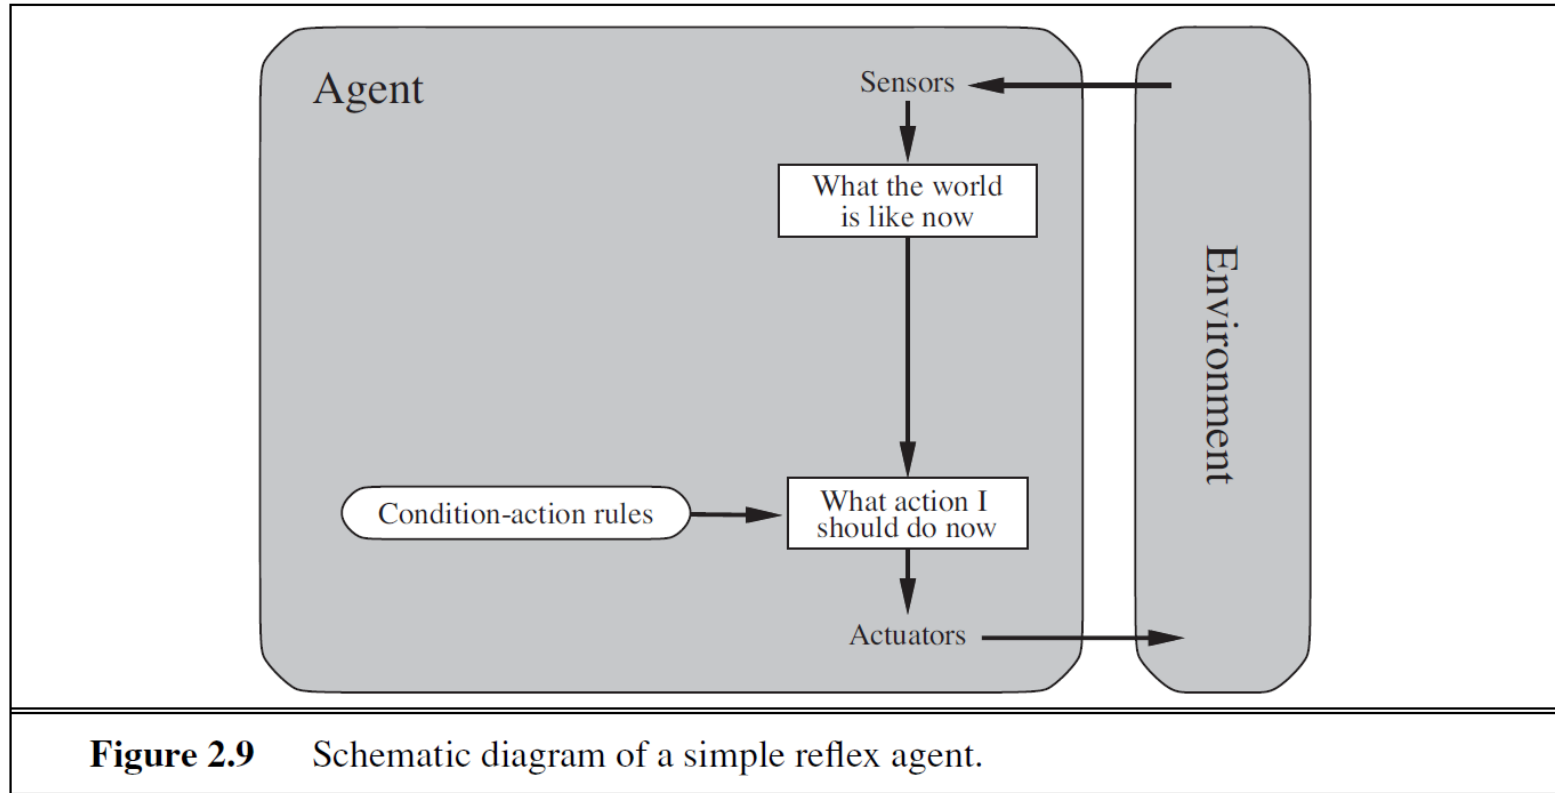
Table-lookup agent

- Drawbacks:
 - Huge table
 - Take a long time to build the table
 - No autonomy
 - Even with learning, need a long time to learn the table entries

Agent types

- Four basic types in order of increasing generality:
 - Simple reflex agents
 - Model-based reflex agents
 - Goal-based agents
 - Utility-based agents
- All can be turned into learning agents

Simple reflex agents



function REFLEX-VACUUM-AGENT(*[location,status]*) returns an action

if *status* = *Dirty* then return *Suck*
else if *location* = *A* then return *Right*
else if *location* = *B* then return *Left*

```
(setq joe (make-agent :name 'joe :body (make-agent-body)
                      :program (make-reflex-vacuum-agent-program)))
```

```
(defun make-reflex-vacuum-agent-program ()
  #'(lambda (percept)
      (let ((location (first percept)) (status (second percept)))
        (cond ((eq status 'dirty) 'Suck)
              ((eq location 'A) 'Right)
              ((eq location 'B) 'Left))))))
```

Model-based reflex agents

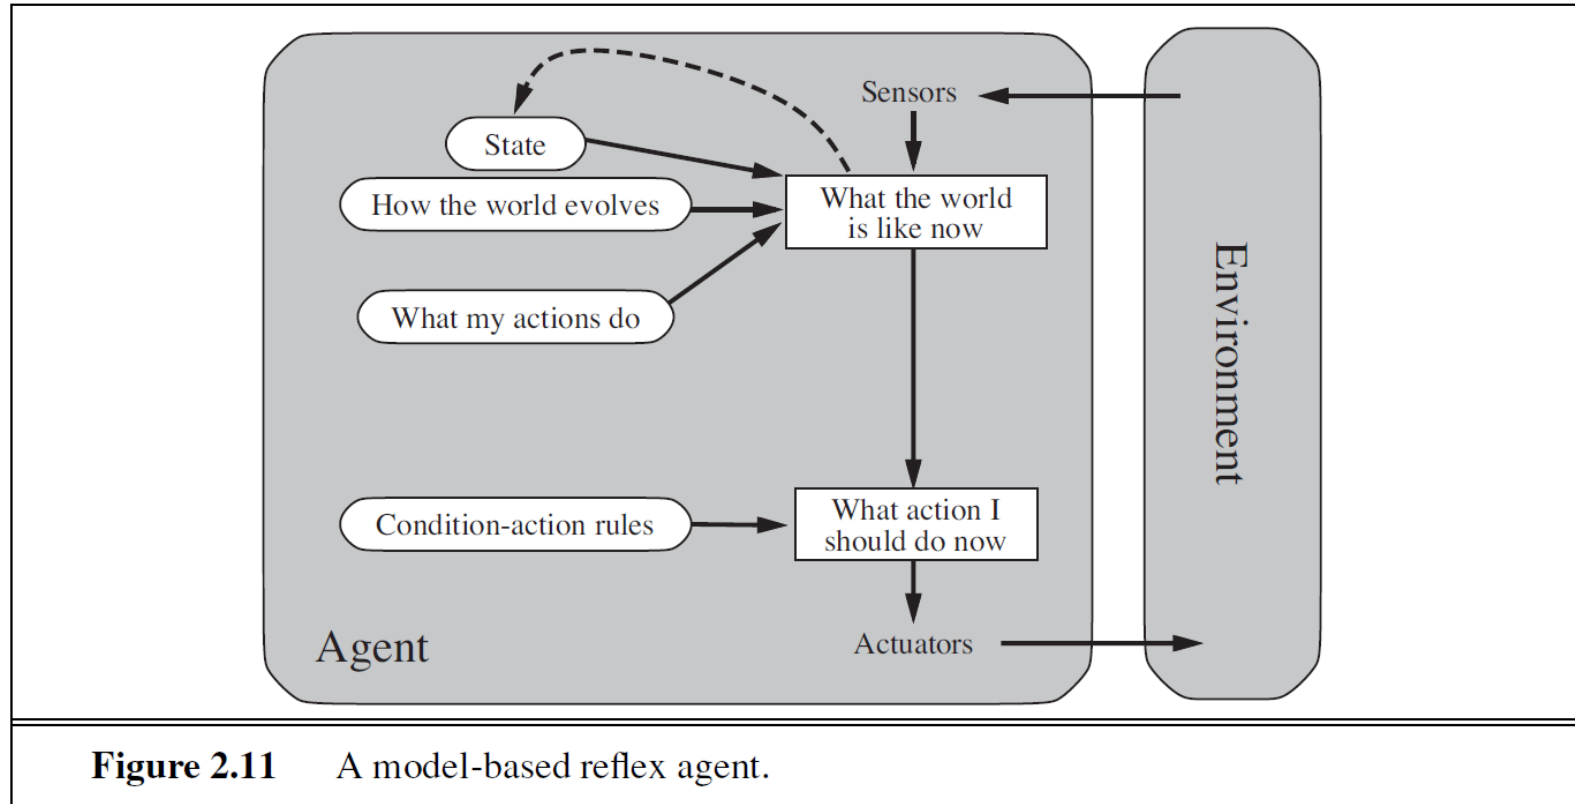


Figure 2.11 A model-based reflex agent.

function REFLEX-VACUUM-AGENT(*[location, status]*) returns an action
static: *last_A*, *last_B*, numbers, initially ∞
if *status* = *Dirty* then ...

```
(defun make-reflex-vacuum-agent-with-state-program ()  
  (let ((last-A infinity) (last-B infinity))  
    #'(lambda (percept)  
      (let ((location (first percept)) (status (second percept)))  
        (incf last-A) (incf last-B)  
        (cond  
          ((eq status 'dirty)  
           (if (eq location 'A) (setq last-A 0) (setq last-B 0))  
           'Suck)  
          ((eq location 'A) (if (> last-B 3) 'Right 'NoOp))  
          ((eq location 'B) (if (> last-A 3) 'Left 'NoOp)))))))
```

Goal-based agents

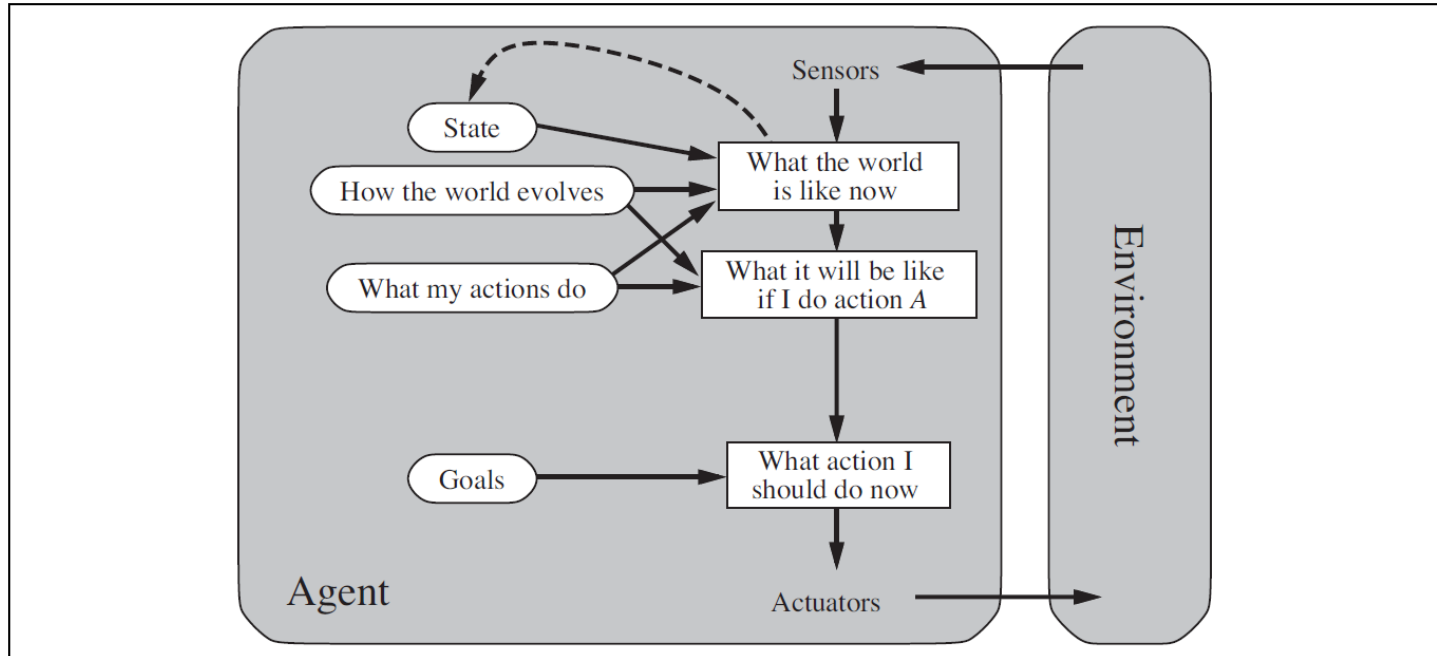


Figure 2.13 A model-based, goal-based agent. It keeps track of the world state as well as a set of goals it is trying to achieve, and chooses an action that will (eventually) lead to the achievement of its goals.

Utility-based agents

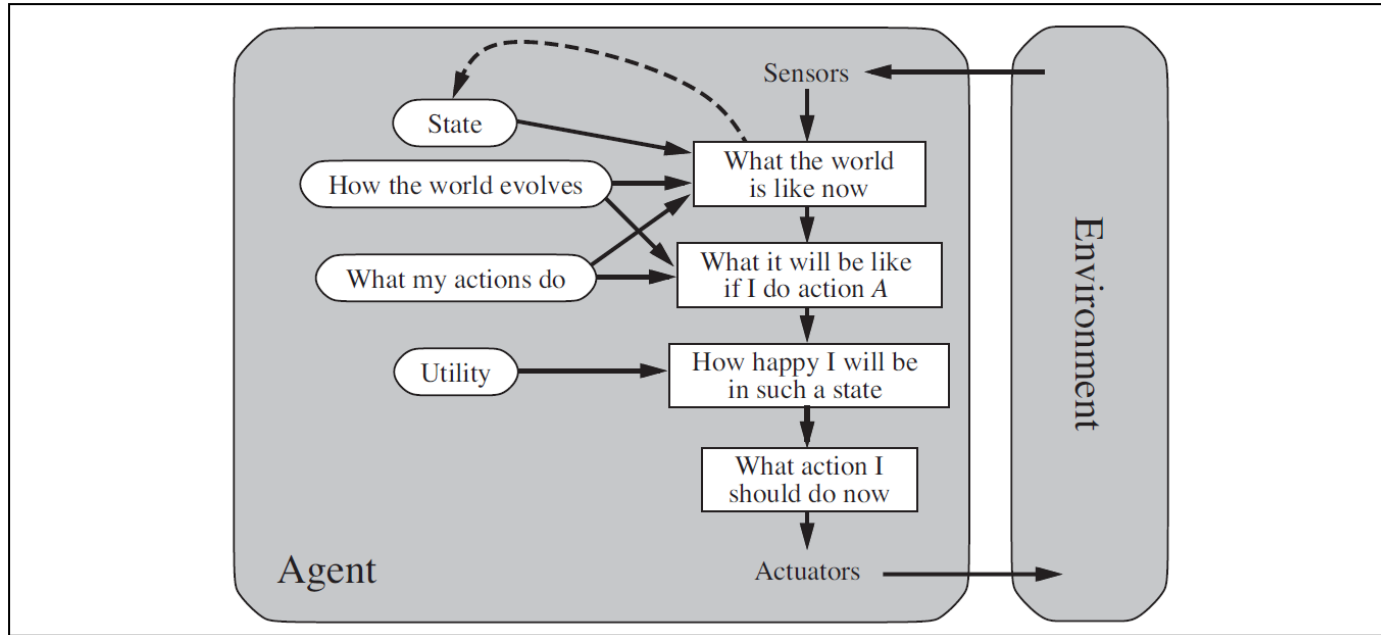


Figure 2.14 A model-based, utility-based agent. It uses a model of the world, along with a utility function that measures its preferences among states of the world. Then it chooses the action that leads to the best expected utility, where expected utility is computed by averaging over all possible outcome states, weighted by the probability of the outcome.

Learning agents

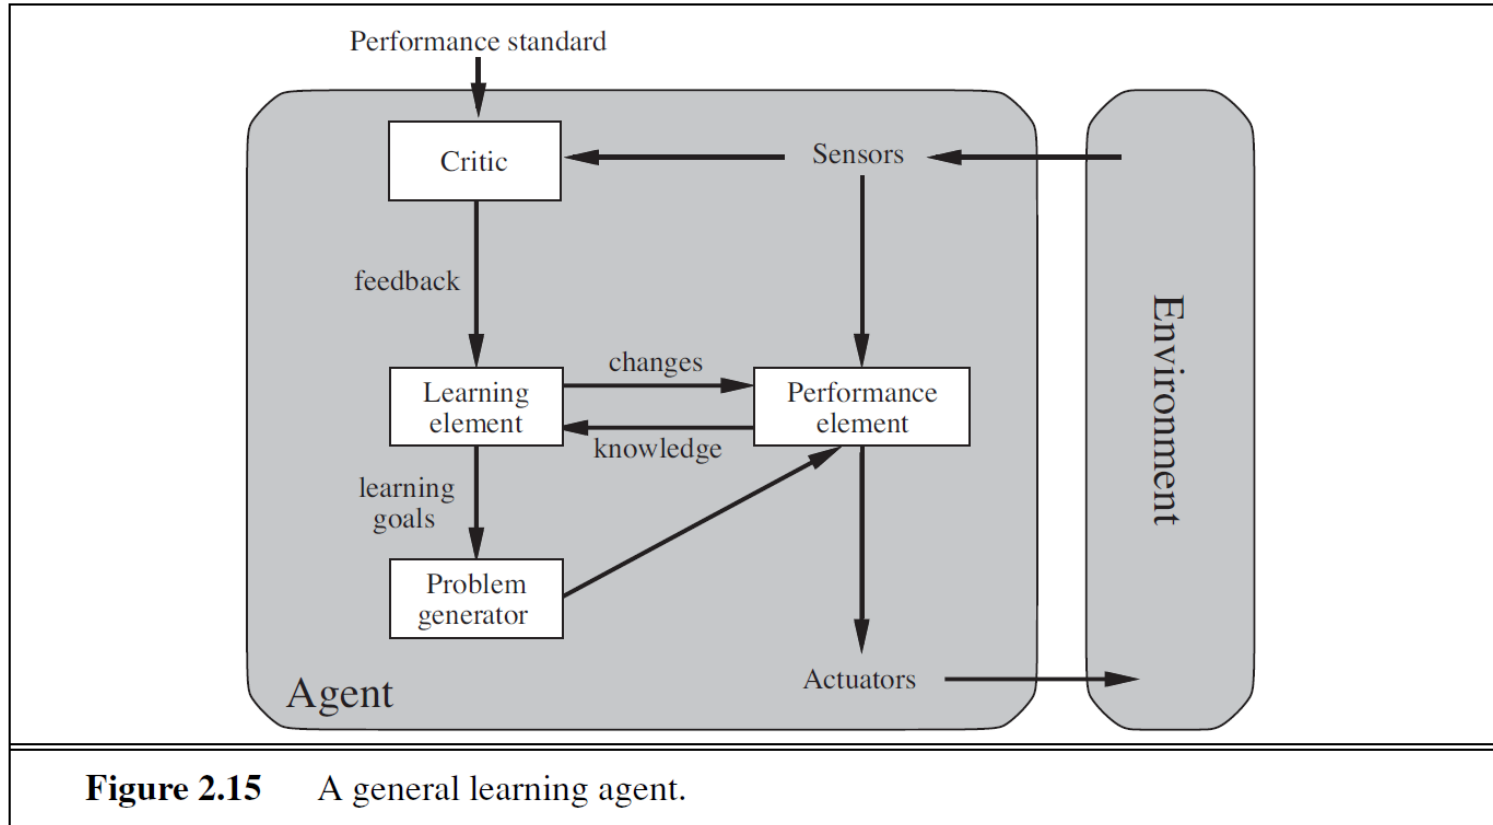


Figure 2.15 A general learning agent.

Summary

Agents interact with environments through actuators and sensors

The agent function describes what the agent does in all circumstances

The performance measure evaluates the environment sequence

A perfectly rational agent maximizes expected performance

Agent programs implement (some) agent functions

PEAS descriptions define task environments

Environments are categorized along several dimensions:

observable? deterministic? episodic? static? discrete? single-agent?

Several basic agent architectures exist:

reflex, reflex with state, goal-based, utility-based

AI