

AI

Artificial Intelligence

8.2.1

Problem Solving and Search (Ch. 3)

Outline

- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms

Cops and Robbers

- Three cops and three robbers come to a river and find a boat that holds two. If the robbers ever outnumber the cops on either bank, the robbers will overpower the cops and escape.
- What is the strategy to cross the river without letting the robbers escape?

Cops and Robbers

- State representation?
- Start State?
- End State?
- Transitions?

Problem-solving agents

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
          state, some description of the current world state
          goal, a goal, initially null
          problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
  action ← RECOMMENDATION(seq, state)
  seq ← REMAINDER(seq, state)
  return action
```



ROMANIA

Neighboring countries: HUNGARY, UKRAINE, MOLDOVA, BULGARIA

Black Sea

Counties and their capital cities:

- SATU MARE: Satu Mare
- MARAMURES: Baia Mare
- BISTRITA-NASAUD: Bistrita
- SUCEAVA: Suceava
- BOTOSANI: Botosani
- IASI: Iasi
- NEAMT: Piatra Neamt
- BACAU: Bacau
- VASLUI: Vaslui
- COVASNA: Sfintu Gheorghe
- VRANCEA: Focsani
- GALATI: Galati
- TULCEA: Tulcea
- CONSTANTA: Constanta
- CELANESE: Galati
- MEHEDINTI: Drobeta-Turnu Severin
- GORJ: Targu Jiu
- VILCEA: Rimnicu Vilcea
- ARGES: Pitesti
- DIMBOVITA: Ploiesti
- PRAHOVA: Ploiesti
- BUZAU: Buzau
- BRASOV: Brasov
- SIBIU: Sibiu
- HUNEDOARA: Deva
- TIMIS: Timisoara
- ARAD: Arad
- BIHOR: Oradea
- CLUJ: Cluj
- MURES: Mures
- HARGHITA: Miercurea Ciuc
- ALBA: Alba Iulia
- MEHEDINTI: Drobeta-Turnu Severin
- DOLJ: Craiova
- TELEORMAN: Alexandria
- GIURGIU: Giurgiu
- BUCURESTI: Bucharest
- IALOMITA: Slobozia
- CALARASI: Calarasi
- BRAILA: Braila

www.mapsofopen.com

Example: Romania

- On holiday in Romania; currently in Arad.
- Flight leaves tomorrow from Bucharest
- Formulate goal:
 - be in Bucharest
- Formulate problem:
 - **states**: various cities
 - **actions**: drive between cities
- Find solution:
 - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

Example: Romania

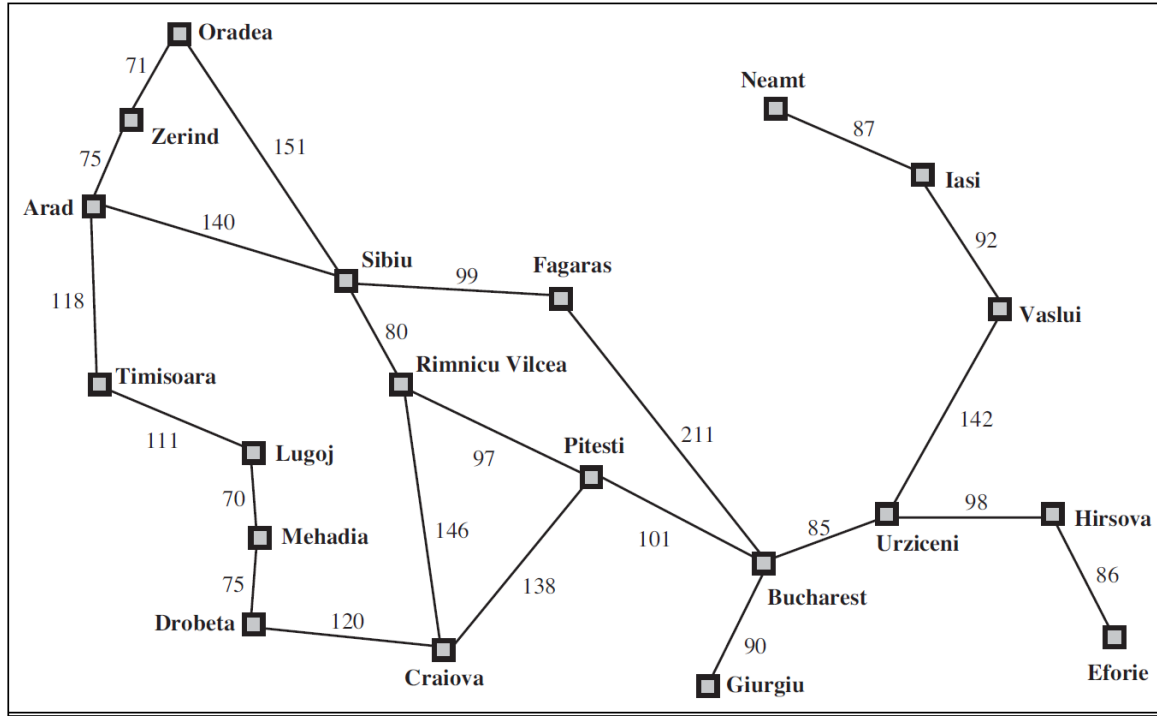


Figure 3.2 A simplified road map of part of Romania.


Problem Types

- Deterministic, fully observable → single-state problem
 - Agent knows exactly which state it will be in; solution is a sequence
- Non-observable → sensorless problem (conformant problem)
 - Agent may have no idea where it is; solution is a sequence
- Nondeterministic and/or partially observable → contingency problem
 - percepts provide new information about current state
 - often interleave search, execution
- Unknown state space → exploration problem

Example: Vacuum World

- Single-state, start in #5.


Solution?



Example: Vacuum World


- Single-state, start in #5.
Solution? [*Right, Suck*]

- Sensorless, start in
 $\{1,2,3,4,5,6,7,8\}$ e.g.,
Right goes to $\{2,4,6,8\}$
Solution?



Example: Vacuum World

- Sensorless, start in $\{1,2,3,4,5,6,7,8\}$ e.g.,
Right goes to $\{2,4,6,8\}$
Solution?
[Right, Suck, Left, Suck]
- Contingency
 - Nondeterministic: *Suck* may dirty a clean carpet
 - Partially observable: location, dirt at current location.
 - Percept: $[L, \text{Clean}]$, i.e., start in #5 or #7
Solution?




Example: Vacuum World

- Sensorless, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}

Solution?

[Right, Suck, Left, Suck]

- Contingency
 - Nondeterministic: *Suck* may dirty a clean carpet
 - Partially observable: location, dirt at current location.
 - Percept: *[L, Clean]*, i.e., start in #5 or #7
- Solution? *[Right, if dirt then Suck]*



Single-state problem formulation

A **problem** is defined by four items:

1. **initial state** e.g., "at Arad"
 2. **actions** or **successor function** $S(x)$ = set of action–state pairs
 - e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
 3. **goal test**, can be
 - **explicit**, e.g., $x = \text{"at Bucharest"}$
 - **implicit**, e.g., $\text{Checkmate}(x)$
 4. **path cost** (additive)
 - e.g., sum of distances, number of actions executed, etc.
 - $c(x,a,y)$ is the step cost, assumed to be ≥ 0
- A **solution** is a sequence of actions leading from the initial state to a goal state

Selecting a State Space

- Real world is absurdly complex
 - state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
 - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, **any** real state "in Arad" must get to **some** real state "in Zerind"
- (Abstract) solution = Set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem

Reflex Agent Solution

- Precompute all paths from any start state to the final state
- Store these paths in a table
- When faced with a new instance of the problem, look up the answer in the table

Cops and Robbers

- Three cops and three robbers come to a river and find a boat that holds two. If the robbers ever outnumber the cops on either bank, the robbers will overpower the cops and escape.
- What is the strategy to cross the river without letting the robbers escape?

Cops and Robbers

- State representation?
- Start State?
- End State?
- Transitions?

Cops and Robbers

- State representation?
 - Number of cops on the source bank
 - Number of robbers on the source bank
 - Number of boats on the source bank
- Start State?
- End State?
- Transitions?
- How many states total?
- Some states are forbidden
- Some states are unreachable

Cops and Robbers

- Search space





Image from Gerhard Wickler

Cops and Robbers


- Solution $331 \rightarrow 310 \rightarrow 321 \rightarrow 300 \rightarrow 311 \rightarrow 110 \rightarrow 221 \rightarrow 020 \rightarrow 031 \rightarrow 010 \rightarrow 021 \rightarrow 000$.
- Is a reflex agent the right method for this problem?
- Idea: separate algorithm/model from problem description
- Idea: planning agent

Vacuum world state space graph



- states?
- actions?
- goal test?
- path cost?

State Space Graph



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

Example: The 8-puzzle

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

- states?
- actions?
- goal test?
- path cost?

Example: The 8-puzzle

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |


Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move
- [Note: optimal solution of n -Puzzle family is NP-hard]

Example: Robotic Assembly




- states?: real-valued coordinates of robot joint angles parts of the object to be assembled
- actions?: continuous motions of robot joints
- goal test?: complete assembly
- path cost?: time to execute

Tree Search Algorithms


- Basic idea:
 - offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. ~expanding states)

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```


Tree search example



Tree search example



Tree search example




Implementation: general tree search

```
function TREE-SEARCH( problem, fringe) returns a solution, or failure  
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)  
  loop do  
    if fringe is empty then return failure  
    node  $\leftarrow$  REMOVE-FRONT(fringe)  
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)  
    fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
```

```
function EXPAND( node, problem) returns a set of nodes  
  successors  $\leftarrow$  the empty set  
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do  
    s  $\leftarrow$  a new NODE  
    PARENT-NODE[s]  $\leftarrow$  node; ACTION[s]  $\leftarrow$  action; STATE[s]  $\leftarrow$  result  
    PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(node, action, s)  
    DEPTH[s]  $\leftarrow$  DEPTH[node] + 1  
    add s to successors  
  return successors
```

Implementation: states vs. nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree includes **state**, **parent node**, **action**, **path cost $g(x)$** , **depth**



- The **Expand** function creates new nodes, filling in the various fields and using the **SuccessorFn** of the problem to create the corresponding states.

Search Strategies


- A search strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
 - completeness: does it always find a solution if one exists?
 - time complexity: number of nodes generated
 - space complexity: maximum number of nodes in memory
 - optimality: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - b : maximum branching factor of the search tree
 - d : depth of the least-cost solution
 - m : maximum depth of the state space (may be ∞)

Uninformed search strategies

- **Uninformed** search strategies use only the information available in the problem definition
 - Breadth-first search
 - Uniform-cost search
 - Depth-first search
 - Depth-limited search
 - Iterative deepening search


Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end




Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end




Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end



Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end



Properties of breadth-first search


- Complete? Yes (if b is finite)
- Time? $1+b+b^2+b^3+\dots +b^d + b(b^d-1) = O(b^{d+1})$
- Space? $O(b^{d+1})$ (keeps every node in memory)
- Optimal? Yes (if cost = 1 per step)
- **Space** is the bigger problem (more than time)

Uniform-cost search

- Expand least-cost unexpanded node
- Implementation:
 - *fringe* = queue ordered by path cost
- Equivalent to breadth-first if step costs all equal
- Complete? Yes, if step cost $\geq \epsilon$
- Time? # of nodes with $g \leq$ cost of optimal solution, $O(b^{\text{ceiling}(C^*/\epsilon)})$ where C^* is the cost of the optimal solution
- Space? # of nodes with $g \leq$ cost of optimal solution, $O(b^{\text{ceiling}(C^*/\epsilon)})$
- Optimal? Yes – nodes expanded in increasing order of $g(n)$


Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front



Properties of depth-first search

- Complete? No: fails in infinite-depth spaces, spaces with loops
 - Modify to avoid repeated states along path
 - complete in finite spaces
- Time? $O(b^m)$: terrible if m is much larger than d
 - but if solutions are dense, may be much faster than breadth-first
- Space? $O(bm)$, i.e., linear space!
- Optimal? No

Demos

- <https://www.cs.usfca.edu/~galles/visualization/DFS.html>
- <https://www.cs.usfca.edu/~galles/visualization/BFS.html>
- <https://www.cs.usfca.edu/~galles/visualization/Algorithms.html>
- <http://www.redblobgames.com/pathfinding/a-star/introduction.html>
- <http://www.greenfoot.org/scenarios/4900>

Depth-limited search

= depth-first search with depth limit l ,
i.e., nodes at depth l have no successors

- Recursive implementation:

```
function DEPTH-LIMITED-SEARCH( problem, limit) returns soln/fail/cutoff
    RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
    cutoff-occurred?  $\leftarrow$  false
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    else if DEPTH[node] = limit then return cutoff
    else for each successor in EXPAND(node, problem) do
        result  $\leftarrow$  RECURSIVE-DLS(successor, problem, limit)
        if result = cutoff then cutoff-occurred?  $\leftarrow$  true
        else if result  $\neq$  failure then return result
    if cutoff-occurred? then return cutoff else return failure
```

Iterative deepening search

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or fail-  
ure  
  inputs: problem, a problem  
  for depth  $\leftarrow$  0 to  $\infty$  do  
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)  
    if result  $\neq$  cutoff then return result
```


Iterative deepening search $l = 0$

Limit = 0









Iterative deepening search $l = 1$

Limit = 1




Iterative deepening search $l = 2$

Limit = 2



Iterative deepening search $l = 3$

Limit = 3



Iterative deepening search

- Number of nodes generated in a depth-limited search to depth d with branching factor b :

$$N_{DLS} = b^0 + b^1 + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$

- Number of nodes generated in an iterative deepening search to depth d with branching factor b :

$$N_{IDS} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

- For $b = 10$, $d = 5$,
 - $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$
 - $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$
- Overhead = $(123,456 - 111,111)/111,111 = 11\%$

Properties of iterative deepening search



- Complete? Yes
- Time? $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space? $O(bd)$
- Optimal? Yes, if step cost = 1

Summary of algorithms

| Criterion | Breadth-First | Uniform-Cost | Depth-First | Depth-Limited | Iterative Deepening |
|-----------|---------------|----------------------------------|-------------|--------------------|---------------------|
| Complete? | Yes* | Yes* | No | Yes, if $l \geq d$ | Yes |
| Time | b^{d+1} | $b^{\lceil C^*/\epsilon \rceil}$ | b^m | b^l | b^d |
| Space | b^{d+1} | $b^{\lceil C^*/\epsilon \rceil}$ | bm | bl | bd |
| Optimal? | Yes* | Yes | No | No | Yes* |

Repeated states

- Failure to detect repeated states can turn a linear problem into an exponential one!



Dynamic Programming

- Split the problem of computing the total cost from Start state to End state into:
 - Past cost to current state S_i
 - Future cost from S_i to end state
- The future cost is equal to 0 at the final state
- In other states:
 - Consider all possible actions a_j
 - Pick the one that minimizes
 - $\text{Cost}(S_i, a_j) + \text{Future Cost}(\text{successor}(S_i, a_j))$

Dynamic Programming

- How to avoid exponential costs?
 - By storing all relevant history in the description of the current state
- Romania example:
 - The future cost of a node doesn't depend on the path that got us there
 - In this case, the current node only needs to store the best cost that got us there
 - DP rule: if the best past cost has been computed for a node, don't compute it again

Dynamic Programming Demos

- <https://www.cs.usfca.edu/~galles/visualization/DPFib.html>
- <https://www.cs.usfca.edu/~galles/visualization/DPLCS.html>
- https://blast.ncbi.nlm.nih.gov/Blast.cgi?PAGE_TYPE=BlastSearch&PROG_DEF=blastn&BLAST_PROG_DEF=blastn&BLAST_SPEC=GlobalAln&LINK_LOC=BlastHomeLink
- <http://experiments.mostafa.io/public/needleman-wunsch/>
- <https://zhanglab.ccmb.med.umich.edu/NW-align/>
- <https://zhanglab.ccmb.med.umich.edu/FASTA/>
- <http://www.cs.yale.edu/homes/radev/nlpclass/slides2017/313.pdf>

Sample Protein Sequence

>gi|186681228|ref|YP_001864424.1| phycoerythrobilin:ferredoxin oxidoreductase

MNSERSDVTLYQPFLDYAIAYMRSRLDLEPYPIPTGFESNSAVVGKGNQEEVVTTSYAFQTAKLRQIRA
AHVQGGNSLQVLNFVIFPHLNYDLPPFFGADLVTLPGGHLIALDMQPLFRDDSAQAKYTEPILPIFHAHQ
QHLSWGGDFPEEAQPFSPAFLWTRPQETAVVETQVFAAFKDYLKAYLDFVEQAEAVTDSQNLVAIKQAQ
LRYLRYRAEKDPARGMFKRFYGAEWTEEYIHGFLFDLERKLTVVK

Nucleic acid codes:

A --> adenosine M --> A C (amino) C --> cytidine S --> G C (strong) G --> guanine W --> A T (weak) T --> thymidine B --
> G T C U --> uridine D --> G A T R --> G A (purine) H --> A C T Y --> T C (pyrimidine) V --> G C A K --> G T (keto) N --
> A G C T (any)

Amino acid codes:

A ALA alanine P PRO proline B ASX aspartate or asparagine Q GLN glutamine C CYS cystine R ARG arginine
D ASP aspartate S SER serine E GLU glutamate T THR threonine F PHE phenylalanine U selenocysteine
G GLY glycine V VAL valine H HIS histidine W TRP tryptophan I ILE isoleucine Y TYR tyrosine K LYS lysine
Z GLX glutamate or glutamine L LEU leucine X any M MET methionine * translation stop N ASN asparagine
- gap of indeterminate length

Uniform Cost Search

- Dynamic programming only works when the graph is acyclic
- The assumption is that we will compute the past cost of a state before computing the past cost of its successors.
- This doesn't work for general graphs


Uniform Cost Search

- The idea of uniform cost search is to order the states by past cost
- Not the same as Dijkstra's algorithm because the graph may not be fully specified at search time
- Furthermore, UCS is intended to find the best cost path only from the start state to the goal state, not between all pairs of states.

Uniform Cost Search

- Explored states – cost known
- Frontier states – known states but cost is unknown
- Unexplored states – the rest

Dijkstra Algorithm



https://en.wikipedia.org/wiki/Dijkstra%27s_algorithm#/media/File:Dijkstra_Animation.gif

Graph Search

function GRAPH-SEARCH(*problem*, *fringe*) **returns** a solution, or failure

closed ← an empty set

fringe ← INSERT(MAKE-NODE(INITIAL-STATE[*problem*]), *fringe*)

loop do

if *fringe* is empty **then return** failure

node ← REMOVE-FRONT(*fringe*)

if GOAL-TEST[*problem*](STATE[*node*]) **then return** SOLUTION(*node*)

if STATE[*node*] is not in *closed* **then**

 add STATE[*node*] to *closed*

fringe ← INSERTALL(EXPAND(*node*, *problem*), *fringe*)

Dijkstra Demos

- <https://www.cs.usfca.edu/~galles/visualization/Dijkstra.html>
- <https://qiao.github.io/PathFinding.js/visual/>
- <https://visualgo.net/de/sssp?slide=1>

Learning the Costs

- The solution depends on the costs
- If the costs are given, fine
 - What if they are not?
- This is a learning problem (the inverse of search)
 - Observed sequence of actions \rightarrow costs
- Sample algorithm
 - Structured perceptron (later)

Summary

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed search strategies
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms

AI